

1

Balanced Trees, Graphs

CSC 340

April 20, 2016

2

Overview

- ✦ **Balanced Search Trees**

- ✦ **Graphs**

3

Balanced Search Trees

- ✦ The efficiency of the binary search tree implementation that we've seen (and implemented) is related to the tree's height. Why?

Sensitive to order of insertion and deletion

4

First: 10, 20, 30, 40, 50, 60, 70
 Second: 40, 20, 60, 10, 30, 50, 70

Balanced Search Trees

- ✦ Consider the following trees that could be generated from the same set of items, based on insertion order



5

Balanced Search Trees

- ✦ Height of a binary search tree of n items:
 - ✦ Maximum: n
 - ✦ Minimum: $\log_2(n + 1)$

6

Balanced Search Trees


- ✦ Search trees can retain their balance despite order of insertions and deletions
 - ✦ 2-3 trees
 - ✦ 2-3-4 trees
 - ✦ Red-black trees
 - ✦ AVL trees

2-3 Trees

- ❖ Have 2-nodes and 3-nodes
 - ❖ A 2-node has one data item and two children
 - ❖ A 3-node has two data items and three children
- ❖ Are general trees, not binary trees
- ❖ Are never taller than a minimum height binary tree
 - ❖ A 2-3 tree with n nodes never has height greater than $\log_2(n+1)$

2-3 Trees - Examples

- ❖ Balanced binary search tree and a 2-3 tree with the same elements



2-3 Trees

- ❖ A leaf may contain either one or two data items
- ❖ To traverse a 2-3 tree, perform the analogue of an in-order traversal
- ❖ Searching a 2-3 tree is as efficient as searching the shortest binary search tree - $O(\log_2 n)$
- ❖ Insertion into a 2-node leaf is simple
- ❖ Insertion into a 3-node leaf causes it to divide

10

2-3 Trees - Insertions

- ❖ A 2-node contains a single data item whose search key S satisfies the following:
 - ❖ $S >$ the left child's search key(s)
 - ❖ $S <$ the right child's search key(s)

11

2-3 Trees - Insertions

- ❖ A 3-node contains two data items whose search keys S and L satisfy the following:
 - ❖ $S >$ the left child's search key(s)
 - ❖ $S <$ the middle child's search key(s)
 - ❖ $L >$ the middle child's search key(s)
 - ❖ $L <$ the right child's search key(s)

12

2-3 Trees - Insertions

- ❖ Locate the leaf at which the search for the item would terminate
- ❖ Insert the new item into the leaf
- ❖ If the leaf now contains only two items, you are done
- ❖ If the leaf now contains three items, $a < b < c$, split the leaf into two nodes containing a and c , and promote b to the parent
- ❖ When the root contains three items
 - ❖ Split the root into two nodes
 - ❖ Create a new root node
 - ❖ Tree grows in height

Deletions are slightly more complicated than I want to go into, but I encourage you to read about it

2-3 Trees

- ✦ Maintaining the balance of a 2-3 tree is relatively easy (not so with the binary search tree)
- ✦ A 2-3 implementation of a table is $O(\log_2 n)$ for all table operations (insert, remove, lookup)
- ✦ A 2-3 tree is a compromise
 - ✦ Searching a 2-3 tree is not more efficient than searching a binary search tree of minimum height
 - ✦ The 2-3 tree might be shorter, but the advantage is offset by the extra comparisons in a node having two values

2-3-4 Trees

- ✦ Have 2-nodes, 3-nodes, and 4-nodes
 - ✦ 4-node has three data items and four children
- ✦ Are general trees, not binary trees
- ✦ Are never taller than a 2-3 tree
- ✦ Search and traversal algorithms are simple extensions of the corresponding 2-3 tree algorithms

Overview

- ✦ Balanced Search Trees
- ✦ **Graphs**

16

Definition

- ❖ A graph G consists of two sets
 - ❖ A set, V , of vertices (or nodes)
 - ❖ A set, E , of edges
- ❖ $G = \{ V, E \}$
- ❖ A subgraph consists of a subset of a graph's vertices and a subset of its edges

17


Terminology

- ❖ **Adjacent vertices:** Two vertices that are joined by an edge
- ❖ **Path:** A sequence of edges that begins at one vertex and ends at another vertex
 - ❖ May pass through the same vertex more than once
- ❖ **Simple Path:** A path that passes through a vertex only once
- ❖ **Cycle:** A path that begins and ends at the same vertex
- ❖ **Simple Cycle:** A cycle that does not pass through a vertex more than once

18

Terminology


- ❖ **Connected Graph:** A graph that has a path between each pair of distinct vertices



19

Terminology


- ✦ Disconnected Graph: A graph that has at least one pair of vertices without a path between them



20

Terminology


- ✦ Complete Graph: A graph that has an edge between each pair of distinct vertices



21

Weighted Graph


- ✦ A graph whose edges have numeric labels



Directed Graph

- ✦ Examples we've seen so far are *undirected*
- ✦ In a directed graph,
 - ✦ Each edge has a direction (directed edges)
 - ✦ Can have two edges between a pair of vertices (one in each direction)
 - ✦ Directed path is a sequence of directed edges between two vertices
 - ✦ Vertex Y is adjacent to vertex X if there is a directed edge from X to Y

Directed Graph Example



Note this is a directed, unweighted graph, but distinct edges can have weights as well, and the weights can vary in each direction

Graphs as ADTs

- ✦ Variations of an ADT graph are possible
 - ✦ Vertices may or may not contain values (city names for example) - many problems are solved with only the relationship among vertices
 - ✦ Directed or undirected
 - ✦ Weighted or unweighted
- ✦ Insertion and deletion operations for graphs apply to vertices and edges
- ✦ Graphs can have traversal operations


Implementations

- ❖ Most common implementation of a graph
 - ❖ Adjacency Matrix
 - ❖ Adjacency List
- ❖ Adjacency matrix for a graph that has N vertices numbers $0, 1, \dots, N - 1$ is an N by N array matrix such that $\text{matrix}[i][j]$ indicates whether an edge exists from vertex i to vertex j (and possibly a weight)

Adjacency Matrix

- ❖ For an unweighted graph, $\text{matrix}[i][j]$ is
 - ❖ 1 (or true) if an edge exists from vertex i to vertex j
 - ❖ 0 (or false) if no edge exists from vertex i to vertex j
- ❖ For a weighted graph, $\text{matrix}[i][j]$ is
 - ❖ The weight of the edge from vertex i to vertex j
 - ❖ Infinity if no edge exists from vertex i to vertex j

Adjacency Matrix



	0	1	2	3
	A	B	C	D
0	A	8	∞	6
1	B	8	∞	9
2	C	∞	9	∞
3	D	6	∞	∞

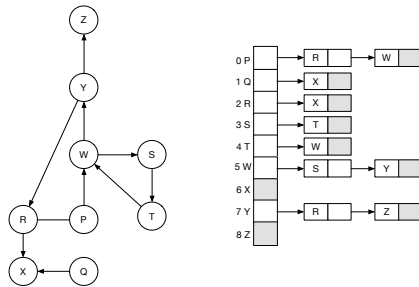
Adjacency Lists

- ✦ Adjacency list for a directed graph that has N vertices numbered $0, 1, \dots, N-1$
 - ✦ An array of N linked lists
 - ✦ The i th linked list has a node for vertex j iff an edge exists from vertex i to vertex j
 - ✦ The list's node can contain either
 - ✦ Vertex j 's value (if any)
 - ✦ An indication of vertex j 's identity

The ints are indices with corresponding vertex value

Adjacency Lists

- ✦ For an undirected graph, treat each edge as if it were two directed edges in opposite directions



Implementation v. Efficiency

- ✦ Two common graph operations:
 - ✦ Determine whether there is an edge from vertex i to vertex j
 - ✦ Find all vertices adjacent to a given vertex i
- ✦ Adjacency Matrix
 - ✦ Supports the first operation more efficiently
- ✦ Adjacency List
 - ✦ Supports the second operation more efficiently
 - ✦ Requires less space than an adjacency matrix

Breadth First Search

- ❖ Algorithm for traversing or searching tree or graph data structure - visits neighbors **first**
- ❖ Take a starting node, N
- ❖ Uses an array to record vertices that have been visited (and, optionally, the distance from the starting node)
- ❖ Uses a queue to record neighbor nodes that need to be visited
- ❖ From starting node, enqueue neighbors, and visit each in turn, setting distance in visited node iff the node has not yet been visited


Breadth First Search

- ❖ Starting at P

Visited Array

P	0
Q	∞
R	1
S	∞
T	∞
W	1
X	∞
Y	∞
Z	∞

Next to Visit Queue



Dequeued R, Adding R's adjacent vertex X


Breadth First Search

- ❖ Starting at P

Visited Array

P	0
Q	∞
R	1
S	∞
T	∞
W	1
X	2
Y	∞
Z	∞

Next to Visit Queue



34



Note that Y is now the value at $W + 1$ (or the weight, if this were a weighted graph)
 Dequeued W, adding W's adjacent vertices Y and S

Breadth First Search

✦ Starting at P

Visited Array	
P	0
Q	∞
R	1
S	2
T	∞
W	1
X	2
Y	2
Z	∞

Next to Visit Queue



35



Dequeued X, X has no adjacent vertices

Breadth First Search

✦ Starting at P

Visited Array	
P	0
Q	∞
R	1
S	2
T	∞
W	1
X	2
Y	2
Z	∞

Next to Visit Queue



36



Dequeued Y, adding Y's adjacent vertex Z

Breadth First Search

✦ Starting at P

Visited Array	
P	0
Q	∞
R	1
S	2
T	∞
W	1
X	2
Y	2
Z	3

Next to Visit Queue



37

Dequeued S, adding S's adjacent vertex T


Breadth First Search

✦ Starting at P

Visited Array

P	0
Q	∞
R	1
S	2
T	3
W	1
X	2
Y	2
Z	3

Next to Visit Queue



38

Dequeued Z, Z has no adjacent vertices


Breadth First Search

✦ Starting at P

Visited Array

P	0
Q	∞
R	1
S	2
T	3
W	1
X	2
Y	2
Z	3

Next to Visit Queue



39

Dequeued T, adding T's adjacent vertex W


Breadth First Search

✦ Starting at P

Visited Array

P	0
Q	∞
R	1
S	2
T	3
W	1
X	2
Y	2
Z	3

Next to Visit Queue




Breadth First Search

❖ Starting at P

Visited Array

P	0
Q	∞
R	1
S	2
T	3
W	1
X	2
Y	2
Z	3

Next to Visit Queue



40

Dequeued W, do not add W's adjacent vertices S and Y as they have already been visited

Depth First Search

- ❖ Explore as far as possible along one path from a node before visiting additional nodes
- ❖ Same concept as BFS - how would you modify the solution for BFS to realize DFS?

41

Use a Stack instead of a Queue

Finding Paths

- ❖ Either BFS or DFS can be used to determine if a path exists between two vertices
- ❖ To determine if a path exists from I to J
 - ❖ Starting from a node I, if the visited array contains a value for node J, a path exists

42